

Machine Learning with Python

Machine Learning

- ▶ Machine Learning is a discipline involving algorithms designed to find patterns in and make predictions about data.
- ▶ It is nearly ubiquitous in our world today, and used in everything from web searches to financial forecasts to studies of the nature of the Universe.
- ▶ This workshop will cover an introduction to scikit-learn, a python machine learning package, and to the central concepts of Machine Learning.

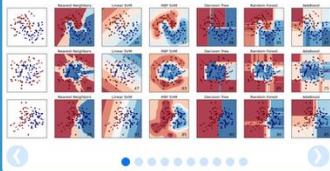
Machine Learning

We will introduce the basic categories of learning problems and how to implement them using scikit-learn.

- ▶ Regression : Predicting Numeric Values
- ▶ Classification : Predicting Categories
- ▶ Clustering : assigning instances to groups.

Getting ready The datasets in scikit-learn are contained within the datasets module. Use the following command to import these datasets:

```
>>> from sklearn import datasets  
>>> import numpy as np
```



scikit-learn

Machine Learning in Python

- Simple and efficient tools for data mining and data analysis
- Accessible to everybody, and reusable in various contexts
- Built on NumPy, SciPy, and matplotlib
- Open source, commercially usable - BSD license

Classification

Identifying to which category an object belongs to.

Applications: Spam detection, Image recognition.

Algorithms: SVM, nearest neighbors, random forest, ... — Examples

Regression

Predicting a continuous-valued attribute associated with an object.

Applications: Drug response, Stock prices.

Algorithms: SVR, ridge regression, Lasso, ... — Examples

Clustering

Automatic grouping of similar objects in

Applications: Customer segmentation, Grouping experiment outcomes

Algorithms: k-Means, spectral clustering, mean-shift, ... — Examples

Dimensionality reduction

Reducing the number of random variables to consider.

Applications: Visualization, Increased efficiency

Algorithms: PCA, feature selection, non-negative matrix factorization. — Examples

Model selection

Comparing, validating and choosing parameters and models.

Goal: Improved accuracy via parameter tuning

Modules: grid search, cross validation, metrics. — Examples

Preprocessing

Feature extraction and normalization.

Application: Transforming input data to text for use with machine learning algorithms

Modules: preprocessing, feature extraction. — Examples

Classification

- ▶ **Description:** Identifying to which category an object belongs to.
- ▶ **Applications:** Spam detection, Image recognition.
- ▶ **Algorithms:** SVM, nearest neighbors, random forest,

Regression

- ▶ **Description:** Predicting a continuous-valued attribute associated with an object.
- ▶ **Applications:** Drug response, Stock prices.
- ▶ **Algorithms:** SVR, ridge regression, Lasso,

Clustering

Automatic grouping of similar objects into sets.

Applications: Customer segmentation,
Grouping experiment outcomes

Algorithms: k-Means, spectral clustering,
mean-shift, ...

Dimensionality Reduction

- ▶ **Description:** Reducing the number of random variables to consider.
- ▶ **Applications:** Visualization, Increased efficiency
- ▶ **Algorithms:** PCA, feature selection, non-negative matrix factorization.

Model selection

- ▶ **Description:** Comparing, validating and choosing parameters and models.
- ▶ **Goal:** Improved accuracy via parameter tuning
- ▶ **Modules:** grid search, cross validation, metrics

Preprocessing

- ▶ **Description:** Feature extraction and normalization.
- ▶ **Application:** Transforming input data such as text for use with machine learning algorithms.
- ▶ **Modules:** preprocessing, feature extraction.